



Module 3

Congestion Control

Approaches towards congestion control

two broad approaches towards congestion control:

end-end congestion control:

- ❖ no explicit feedback from network
- ❖ congestion inferred from end-system observed loss, delay
- ❖ approach taken by TCP


network-assisted congestion control:

- ❖ routers provide feedback to end systems
 - single bit indicating congestion (SNA, DECbit, TCP/IP ECN, ATM)
 - explicit rate for sender to send at

Network Assisted Congestion Control

- Direct feedback may be sent from a network router to the sender. This form of notification typically takes the form of a **choke packet**.
- The second form : when a router marks/updates a field in a packet to indicate congestion.

Upon receipt of a marked packet, the receiver then notifies the sender of the congestion indication.



Two feedback pathways for network-induced congestion information

Case study: ATM ABR congestion control


ABR: available bit rate:

- “elastic service”
- if sender’s path “underloaded” :
 - sender should use available bandwidth
- if sender’s path congested:
 - sender throttled to minimum guaranteed rate

RM (resource management) cells:

- sent by sender, interspersed with data cells
- Each data cell contains an explicit forward congestion indication (EFCI) bit.
- bits in RM cell set by switches (“*network-assisted*”)
 - *NI bit*: no increase in rate (mild congestion)
 - *CI bit*: congestion indication
- RM cells returned to sender by receiver, with bits intact

Case study: ATM ABR congestion control




- two-byte ER (explicit rate) field in RM cell
 - congested switch may lower ER value in cell
 - Sender's send rate thus max supportable rate on path
- EFCI bit in data cells: set to 1 in congested switch
 - if data cell preceding RM cell has EFCI set, receiver sets CI bit in returned RM cell

TCP congestion control: additive increase multiplicative decrease

- ❖ *approach*: sender increases transmission rate (window size), probing for usable bandwidth, until loss occurs
 - *additive increase*: increase **cwnd** by 1 MSS every RTT until loss detected
 - *multiplicative decrease*: cut **cwnd** in half after loss

AIMD saw tooth behavior: probing for bandwidth




Congestion Control Introduction

- If no congestion : increase the rate
 - If congestion : reduces the send rate
 - This approach raises three questions:
 - how does a TCP sender limit the rate at which it sends traffic into its connection
 - how does a TCP sender perceive that there is congestion on the path between itself and the destination
 - what algorithm should the sender use to change its send rate
-

TCP Congestion Control: details

How does a TCP sender limit the rate at which it sends traffic into its connection



- ❖ sender limits transmission:

$$\text{LastByteSent} - \text{LastByteAcked} \leq \text{cwnd}$$

- ❖ the amount of unacknowledged data at a sender may not exceed the minimum of cwnd and rwnd
- ❖ **cwnd** is dynamic, function of perceived network congestion

TCP sending rate:

- ❖ *roughly*: send cwnd bytes, wait RTT for ACKS, then send more bytes

$$\text{rate} \approx \frac{\text{cwnd}}{\text{RTT}} \text{ bytes/sec}$$

Congestion details

- **How a TCP sender perceives that there is congestion on the path between itself and the destination**
 - **“loss event”** at a TCP sender
 - occurrence of either a timeout or
 - the receipt of three duplicate ACKs from the receiver
 - Loss event is the indication of congestion on the sender-to-receiver path
 - When a loss event doesn't occur
 - acknowledgments for previously unacknowledged segments will be received at the TCP sender
-

Congestion Details


- **How should a TCP sender determine the rate at which it should send**
 - Three guiding principles:
 - A lost segment implies congestion, and hence, the TCP sender's rate should be decreased when a segment is lost
 - An acknowledged segment indicates that the network is delivering the sender's segments to the receiver, and hence, the sender's rate can be increased when an ACK arrives for a previously unacknowledged segment.
 - Bandwidth probing :
 - ACKs – increase rate
 - Loss event – decrease rate
-

TCP Congestion Control Algorithm

- First described in [Jacobson 1988]
 - Standardized in [RFC 5681].
 - The algorithm has three major components :
 - (1) **slow start** :
 - value of cwnd begins at 1 MSS and increases by 1 MSS every time a transmitted segment is first acknowledged
 - (2) **congestion avoidance** :
 - value of cwnd equals ssthresh, congestion avoidance mode starts. Tcp increases cwnd more cautiously
 - (3) **fast recovery** :
 - if three duplicate ACKs are detected, in which case TCP performs a fast retransmit
-

TCP Slow Start

- ❖ when connection begins, increase rate exponentially until first loss event:
 - initially `cwnd` = 1 MSS
 - double `cwnd` every RTT
 - done by incrementing `cwnd` for every ACK received
- ❖ *summary*: initial rate is slow but ramps up exponentially fast
- ❖ Exponential growth ends:
 - Loss event occurs
 - `Cwnd` equals `ssthresh`, slow start ends and TCP enters **congestion avoidance**



TCP: detecting, reacting to loss

- ❖ loss indicated by timeout:
 - `cwnd` set to 1 MSS;
 - window then grows exponentially (as in slow start) to threshold, then grows linearly
 - ❖ loss indicated by 3 duplicate ACKs: TCP RENO
 - dup ACKs indicate network capable of delivering some segments
 - `cwnd` is cut in half window then grows linearly
 - ❖ TCP Tahoe always sets `cwnd` to 1 (timeout or 3 duplicate acks)
-


TCP: switching from slow start to CA

Q: when should the exponential increase switch to linear?

A: when **cwnd** gets to 1/2 of its value before timeout.

Implementation:

- ❖ variable **ssthresh**
- ❖ on loss event, **ssthresh** is set to 1/2 of **cwnd** just before loss event



Summary: TCP Congestion Control

